

MK

MORGAN KAUFMANN

COMPUTER ORGANIZATION AND DESIGN  
The Hardware/Software Interface

5<sup>th</sup>

Edition

Chapter 3 & Appendix B  
(continued)

Arithmetic for Computers

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Floating Point

- Representation for non-integral numbers
  - Including very small and very large numbers
- Like scientific notation
  - $-2.34 \times 10^{56}$  ← normalized
  - $+0.002 \times 10^{-4}$  ← not normalized
  - $+987.02 \times 10^9$  ← not normalized
- In binary
  - $\pm 1.xxxxxxx_2 \times 2^{yyy}$
- Types float and double in C

\$3.5 Floating Point

Chapter 3 — Arithmetic for Computers — 2

## Floating Point Standard

- Defined by IEEE standard 754-1985
- Developed in response to divergence of representations
  - Portability issues for scientific code
- Now almost universally adopted
- Two representations
  - Single precision (32-bit)
  - Double precision (64-bit)

Assignment Project Exam Help

Chapter 3 — Arithmetic for Computers — 3

<https://powcoder.com>

## IEEE Floating-Point Format

single: 8 bits  
double: 11 bits

single: 23 bits  
double: 52 bits

S	Exponent	Fraction
---	----------	----------

$$x = (-1)^S \times (1 + \text{Fraction}) \times 2^{(\text{Exponent} - \text{Bias})}$$

- S: sign bit (0 ⇒ non-negative, 1 ⇒ negative)
- Normalized significand:  $1.0 \leq |\text{significand}| < 2.0$ 
  - Always has a leading pre-binary-point 1 bit, so no need to represent it explicitly (hidden bit)
  - Significand is Fraction with the "1." restored
- Exponent: actual exponent + Bias
  - Ensures exponent is unsigned
  - Single: Bias = 127; Double: Bias = 1023

Chapter 3 — Arithmetic for Computers — 4

Chapter 3 — Arithmetic for Computers

2

## Single-Precision Range

- Exponents 00000000 and 11111111 reserved
- Smallest value
  - Exponent: 00000001  
⇒ actual exponent = 1 - 127 = -126
  - Fraction: 000...00 ⇒ significand = 1.0
  - $\pm 1.0 \times 2^{-126} \approx \pm 1.2 \times 10^{-38}$
- Largest value
  - exponent: 11111110  
⇒ actual exponent = 254 - 127 = +127
  - Fraction: 111...11 ⇒ significand  $\approx 2.0$
  - $\pm 2.0 \times 2^{+127} \approx \pm 3.4 \times 10^{+38}$

Assignment Project Exam Help

Chapter 3 — Arithmetic for Computers — 5

<https://powcoder.com>

Add WeChat powcoder

## Double-Precision Range

- Exponents 0000...00 and 1111...11 reserved
- Smallest value
  - Exponent: 00000000001  
⇒ actual exponent = 1 - 1023 = -1022
  - Fraction: 000...00 ⇒ significand = 1.0
  - $\pm 1.0 \times 2^{-1022} \approx \pm 2.2 \times 10^{-308}$
- Largest value
  - Exponent: 11111111110  
⇒ actual exponent = 2046 - 1023 = +1023
  - Fraction: 111...11 ⇒ significand  $\approx 2.0$
  - $\pm 2.0 \times 2^{+1023} \approx \pm 1.8 \times 10^{+308}$

Chapter 3 — Arithmetic for Computers — 6

## Floating-Point Precision

- Relative precision
  - all fraction bits are significant
  - Single: approx  $2^{-23}$ 
    - Equivalent to  $23 \times \log_{10} 2 \approx 23 \times 0.3$   
 $\approx 6$  decimal digits of precision
  - Double: approx  $2^{-52}$ 
    - Equivalent to  $52 \times \log_{10} 2 \approx 52 \times 0.3$   
 $\approx 16$  decimal digits of precision

Assignment Project Exam Help

Chapter 3 — Arithmetic for Computers — 7

<https://powcoder.com>

Add WeChat powcoder

## Binary Refresher

- When we look at a number like  $10110_2$ , we're seeing it as:  
$$1(2^4) + 0(2^3) + 1(2^2) + 1(2^1) + 0(2^0)$$
$$= 16 + 4 + 2 = 22_{10}$$

Chapter 1 — Computer Abstractions and Technology — 8

## Binary Decimal Points

- In decimal, 12.63 is the same as
$$1(10^1) + 2(10^0) + 6(10^{-1}) + 3(10^{-2})$$
- In binary, 101.01<sub>2</sub> is the same as
$$1(2^2) + 0(2^1) + 1(2^0) + 0(2^{-1}) + 1(2^{-2})$$
$$= 4 + 1 + 0.25 = 5.25$$

Assignment Project Exam Help

Chapter 1 — Computer Abstractions and Technology — 9

<https://powcoder.com>

## Useful Exponent Identities

- $a^b * a^c = a^{b+c}$ 
  - Why memorize more than  $2^{10}$  when we can just break them down?
$$2^{35} = 2^5 * 2^{30}$$
$$= 2^5 * 2^{10} * 2^{10} * 2^{10} = 32 \text{ MB}$$
- $a^{-b} = \frac{1}{a^b}$ 
  - When we have decimal terms, instead of seeing  $2^{-1}$ ,  $2^{-2}$ , etc. use  $\frac{1}{2^1}$ ,  $\frac{1}{2^2}$ , etc.

Chapter 1 — Computer Abstractions and Technology — 10

## More Binary

- Dividing (shifting right) by  $2_{10}$  is the same as moving the decimal point one place to the left.
  - Same reasoning as when we divide by 10 in base 10
- Multiplying (shifting left) works the same way

Assignment Project Exam Help

<https://powcoder.com>

## Floating-Point Example

- Represent -0.75
  - $-0.75 = (-1)^1 \times 1.1_2 \times 2^{-1}$
  - $S = 1$
  - Fraction =  $1000...00_2$
  - Exponent =  $-1 + \text{Bias}$ 
    - Single:  $-1 + 127 = 126 = 01111110_2$
    - Double:  $-1 + 1023 = 1022 = 0111111110_2$
- Single:  $1011111101000...00$
- Double:  $101111111101000...00$

Chapter 3 — Arithmetic for Computers — 12

## Floating-Point Example

- What number is represented by the single-precision float  
**11000000101000...00**
  - S = 1
  - Fraction = **01000...00**<sub>2</sub>
  - Exponent = **10000001**<sub>2</sub> = 129
- $x = (-1)^1 \times (1 + 0.01_2) \times 2^{(129 - 127)}$   
 $= (-1) \times 1.25 \times 2^2$   
 $= -5.0$

Assignment Project Exam Help

Chapter 3 — Arithmetic for Computers 13

<https://powcoder.com>

## Add WeChat powcoder

### IEEE 754-1985 Specials

- We reserve all 0s and all 1s in the exponent. This is why:
  - 011111110000...00** =  $+\infty$
  - 111111110000...00** =  $-\infty$
  - X1111111[non-zero]** = NaN
    - e.g., square root of a negative number
  - X000000000000...00** = 0
    - ...there's actually a positive zero and a negative zero

Chapter 1 — Computer Abstractions and Technology — 14

## Floating-Point Addition

- Consider a 4-digit decimal example
  - $9.999 \times 10^1 + 1.610 \times 10^{-1}$
- 1. Align decimal points
  - Shift number with smaller exponent
  - $9.999 \times 10^1 + 0.016 \times 10^1$
- 2. Add significands
  - $9.999 \times 10^1 + 0.016 \times 10^1 = 10.015 \times 10^1$
- 3. Normalize result & check for over/underflow
  - $1.0015 \times 10^2$
- 4. Round and renormalize if necessary
  - $1.002 \times 10^2$

Assignment Project Exam Help

Chapter 3 — Arithmetic for Computers 15

<https://powcoder.com>

Add WeChat powcoder

## Floating-Point Addition

- Now consider a 4-digit binary example
  - $1.000_2 \times 2^{-1} + -1.110_2 \times 2^{-2}$  ( $0.5 + -0.4375$ )
- 1. Align binary points
  - Shift number with smaller exponent
  - $1.000_2 \times 2^{-1} + -0.111_2 \times 2^{-1}$
- 2. Add significands
  - $1.000_2 \times 2^{-1} + -0.111_2 \times 2^{-1} = 0.001_2 \times 2^{-1}$
- 3. Normalize result & check for over/underflow
  - $1.000_2 \times 2^{-4}$ , with no over/underflow
- 4. Round and renormalize if necessary
  - $1.000_2 \times 2^{-4}$  (no change) = 0.0625

Chapter 3 — Arithmetic for Computers — 16



## FP Instructions in MIPS

- FP hardware is coprocessor 1
  - Adjunct processor that extends the ISA
- Separate FP registers
  - 32 single-precision: \$f0, \$f1, ... \$f31
  - Paired for double-precision: \$f0/\$f1, \$f2/\$f3, ...
    - Release 2 of MIPS ISA supports 32 × 64-bit FP reg's
- FP instructions operate only on FP registers
  - Programs generally don't do integer ops on FP data, or vice versa
  - More registers with minimal code-size impact
- FP load and store instructions
  - lwc1, ldc1, swc1, sdc1
    - e.g., ldc1 \$f8, 32(\$sp)

Assignment Project Exam Help

Chapter 3 — Arithmetic for Computers 17

<https://powcoder.com>

Add WeChat powcoder

## FP Instructions in MIPS

- Single-precision arithmetic
  - add.s, sub.s, mul.s, div.s
    - e.g., add.s \$f0, \$f1, \$f6
- Double-precision arithmetic
  - add.d, sub.d, mul.d, div.d
    - e.g., mul.d \$f4, \$f4, \$f6
- Single- and double-precision comparison
  - c.xx.s, c.xx.d (xx is eq, lt, le, ...)
  - Sets or clears FP condition-code bit
    - e.g. c.lt.s \$f3, \$f4
- Branch on FP condition code true or false
  - bc1t, bc1f
    - e.g., bc1t TargetLabel

Chapter 3 — Arithmetic for Computers — 18

FP Example: °F to °C

- C code:

```
float f2c (float fahr) {  
    return ((5.0/9.0)*(fahr - 32.0));  
}
```

  - fahr in \$f12, result in \$f0, literals in global memory space
- Compiled MIPS code:

```
f2c: lwc1    $f16, const5($gp)  
     lwc2    $f18, const9($gp)  
     div.s   $f16, $f16, $f18  
     lwc1    $f18, const32($gp)  
     sub.s   $f18, $f12, $f18  
     mul.s   $f0,  $f16, $f18  
     jr      $ra
```

Assignment Project Exam Help

Chapter 3 — Arithmetic for Computers — 19

<https://powcoder.com>

Add WeChat powcoder

FP Example: Array Multiplication

- $X = X + Y \times Z$ 
  - All  $32 \times 32$  matrices, 64-bit double-precision elements
- C code:

```
void mm (double x[][],  
         double y[][], double z[][]) {  
    int i, j, k;  
    for (i = 0; i! = 32; i = i + 1)  
        for (j = 0; j! = 32; j = j + 1)  
            for (k = 0; k! = 32; k = k + 1)  
                x[i][j] = x[i][j]  
                        + y[i][k] * z[k][j];  
}
```
- Addresses of x, y, z in \$a0, \$a1, \$a2, and i, j, k in \$s0, \$s1, \$s2

Chapter 3 — Arithmetic for Computers — 20

FP Example: Array Multiplication

■ MIPS code:

li	\$t1, 32	# \$t1 = 32 (row size/loop end)
li	\$s0, 0	# i = 0; initialize 1st for loop
L1:	li	\$s1, 0 # j = 0; restart 2nd for loop
L2:	li	\$s2, 0 # k = 0; restart 3rd for loop
	sll	\$t2, \$s0, 5 # \$t2 = i * 32 (size of row of x)
	addu	\$t2, \$t2, \$s1 # \$t2 = i * size(row) + j
	sll	\$t2, \$t2, 3 # \$t2 = byte offset of [i][j]
	addu	\$t2, \$a0, \$t2 # \$t2 = byte address of x[i][j]
	l.d	\$f4, 0(\$t2) # \$f4 = 8 bytes of x[i][j]
L3:	sll	\$t0, \$s2, 5 # \$t0 = k * 32 (size of row of z)
	addu	\$t0, \$t0, \$s1 # \$t0 = k * size(row) + j
	sll	\$t0, \$t0, 3 # \$t0 = byte offset of [k][j]
	addu	\$t0, \$a2, \$t0 # \$t0 = byte address of z[k][j]
	l.d	\$f16, 0(\$t0) # \$f16 = 8 bytes of z[k][j]

...

Assignment Project Exam Help

Chapter 3 — Arithmetic for Computers — 21

<https://powcoder.com>

Add WeChat powcoder

FP Example: Array Multiplication

...

sll	\$t0, \$s0, 5	# \$t0 = i*32 (size of row of y)
addu	\$t0, \$t0, \$s2	# \$t0 = i*size(row) + k
sll	\$t0, \$t0, 3	# \$t0 = byte offset of [i][k]
addu	\$t0, \$a1, \$t0	# \$t0 = byte address of y[i][k]
l.d	\$f18, 0(\$t0)	# \$f18 = 8 bytes of y[i][k]
mul.d	\$f16, \$f18, \$f16	# \$f16 = y[i][k] * z[k][j]
add.d	\$f4, \$f4, \$f16	# f4=x[i][j] + y[i][k]*z[k][j]
addiu	\$s2, \$s2, 1	# \$k k + 1
bne	\$s2, \$t1, L3	# if (k != 32) go to L3
s.d	\$f4, 0(\$t2)	# x[i][j] = \$f4
addiu	\$s1, \$s1, 1	# \$j = j + 1
bne	\$s1, \$t1, L2	# if (j != 32) go to L2
addiu	\$s0, \$s0, 1	# \$i = i + 1
bne	\$s0, \$t1, L1	# if (i != 32) go to L1

Chapter 3 — Arithmetic for Computers — 22

## Accurate Arithmetic

- IEEE Std 754 specifies additional rounding control
  - Extra bits of precision (guard, round, sticky)
  - Choice of rounding modes
  - Allows programmer to fine-tune numerical behavior of a computation
- Not all FP units implement all options
  - Most programming languages and FP libraries just use defaults
- Trade-off between hardware complexity, performance, and market requirements

Assignment Project Exam Help

Chapter 3 — Arithmetic for Computers — 23

<https://powcoder.com>

Add WeChat powcoder

## Associativity

- Parallel programs may interleave operations in unexpected orders
    - Assumptions of associativity may fail
- |   |           |          |           |
|---|-----------|----------|-----------|
|   |           | (x+y)+z  | x+(y+z)   |
| x | -1.50E+38 |          | -1.50E+38 |
| y | 1.50E+38  | 0.00E+00 |           |
| z | 1.0       | 1.0      | 1.50E+38  |
|   |           | 1.00E+00 | 0.00E+00  |
- Need to validate parallel programs under varying degrees of parallelism